

What is *output sensitivity*? Design & analyze algorithms for the convex-hull problem

Output sensitivity (informal). An algorithm is *output sensitive* when its running time depends not only on the input size n but also on the size h of the output. For convex hull in the plane, n is the number of input points and h the number of points on the hull. An output-sensitive convex-hull algorithm runs in time $T(n,h)$ that is smaller when h is small.

Algorithms

1. Graham scan

- Idea: sort points by polar angle about a pivot, then scan to build hull using a stack.
- Time: $O(n \log n)$ dominated by sorting. Not output sensitive.

2. Jarvis march (Gift wrapping)

- Idea: start at extreme point, repeatedly find the next hull point by scanning all points for the most counterclockwise point.
- Time: $O(nh)$. Output sensitive — good when h is small (e.g., constant hull size).

3. Chan's algorithm (combines Jarvis + Graham)

- Idea: partition points into groups of size m , compute hulls of groups (Graham), then do a Jarvis-type wrap using tangents to group hulls; choose m and iterate doubling until full hull found.
- Time: $O(n \log h)$. This is output sensitive and optimal for comparison-based models up to constant factors.

4. Kirkpatrick–Seidel (“ultimate planar convex hull”)

- Time: $O(n \log h)$ as well (divide-and-conquer using a “marriage-before-conquest” step).

Design sketch for Chan's algorithm (practical output-sensitive method):

- Pick parameter m (e.g. try $m=2,4,8,\dots$).
- Partition input into $\lceil n/m \rceil$ groups of at most m points.
- Compute each group's convex hull by Graham scan in $O(m \log m)$ per group → total $O((n/m) \cdot m \log m) = O(n \log m)$.
- Use a Jarvis-step that finds the next hull vertex among all group hulls by computing tangent points to each group hull in $O(\log m)$ per group; each step costs $O((n/m) \log m)$.
- If actual $h \leq m$, Jarvis completes in h steps. Choose m doubling until it succeeds. Overall cost gives $O(n \log h)$.

Example (Convex Hull)

Input points:

A(1,1), B(2,5), C(5,4), D(6,1), E(3,3), F(4,2)

Visual hull points = A → B → C → D (4 points).

Output-sensitive runtimes:

- $n = 6, h = 4$
- **Jarvis march:** $O(nh) = 6 \times 4 = 24$ ops
- **Graham scan:** $O(n \log n) = 6 \log 6 \approx 15$ ops
- **Chan's:** $O(n \log h) = 6 \log 4 \approx 12$ ops

If h were only 2 or 3, Jarvis becomes much faster — this shows output sensitivity.

Define NP-complete and the Vertex-Cover problem. Is Vertex-Cover NP-complete? Justify.

Definitions.

- **NP:** class of decision problems for which a “yes” instance has a certificate (witness) verifiable in polynomial time by a deterministic Turing machine.
- **NP-hard:** a problem L is NP-hard if every problem in NP can be polynomial-time reduced to L. Informally: L is at least as hard as every problem in NP.
- **NP-complete:** a problem is NP-complete if it is in NP *and* NP-hard. These are the hardest problems in NP.

Vertex-Cover (decision form).

Input: Graph $G=(V,E)$ and integer k .

Question: Is there a subset $C \subseteq V$ with $|C| \leq k$ such that every edge has at least one endpoint in

C ? (Such a C is a vertex cover.)

Vertex-Cover is NP-complete:

1. **Vertex-Cover ∈ NP.** Certificate: a set C of $\leq k$ vertices. We can check in polynomial time that every edge has an endpoint in C .
2. **Vertex-Cover is NP-hard.** Standard polynomial-time reduction from **Independent Set** (or from CLIQUE) which are known NP-complete:

- Independent Set decision: given G and k , is there an independent set S of size $\geq k$ (no edges among vertices in S)?
- Reduction: Given $((G,k)$ for Independent Set, produce (G,k') for Vertex Cover with $k'=|V|-k$.
- Correctness: S is an independent set of size k iff $V \setminus S$ is a vertex cover of size $|V|-k$. This is a polynomial, simple transformation, so Vertex-Cover is NP-hard.

3. Combined \Rightarrow Vertex-Cover is NP-complete.

Example

Graph:

Edges: $(A,B), (B,C), (C,D), (A,D)$

Vertex cover of size 2 exists?

Try $C = \{B, C\}$:

- (A,B) covered by B
- (B,C) covered by B or C
- (C,D) covered by C
- (A,D) NOT covered \rightarrow FAIL

Try $C = \{A, C\}$:

- (A,B) covered by A
- (B,C) covered by C
- (C,D) covered by C
- (A,D) covered by A \rightarrow SUCCESS

Thus vertex cover size ≤ 2 exists.

NP-hardness example using reduction from Independent Set

Graph: $A-B-C-D$ as above.

Independent set of size 2: $\{A, C\}$.

Vertex cover size = $|V| - |IS| = 4 - 2 = 2 \rightarrow$ matches above.

This shows the reduction works.

Plane-sweep algorithm for line-segment intersection detection in $O(n \log n)$

Problem: Given set S of n segments in the plane, determine whether any pair intersects.

Algorithm (sweep line — Bentley–Ottmann style simplified for detection):

Data structures:

- Event queue Q : all segment endpoints, sorted by x-coordinate (ties broken by y). (Use balanced BST or priority queue.) Size $2n$.
- Status structure T : ordered set (balanced BST) of segments currently intersecting the vertical sweep line, ordered by intersection point's y-coordinate with the sweep line. Supports insert, delete, predecessor, successor in $O(\log n)$.

Procedure:

1. Initialize Q with all endpoints (mark as left or right endpoint for each segment; for vertical segments define ordering consistently).
2. While Q not empty:
 - Extract next event point p (left-to-right). Let the set of segments whose endpoint is p be processed.
 - If p is a left endpoint of segment s : insert s into T ; check s 's immediate predecessor and successor in T for intersection with s . If any intersection found → return “yes”.
 - If p is a right endpoint of segment s : find predecessor u and successor v of s in T . Delete s from T . If both u and v exist, check u vs v for intersection.
 - (If you want to report all intersections, you must also enqueue intersection events when two neighbors swap; for detection-only, neighbor checks suffice.)
3. If finish without finding intersection → return “no”.

Correctness sketch: If two segments intersect, at the x-coordinate of their leftmost intersection the sweep line will have them adjacent in T at some event; the predecessor/successor checks find the intersection. The algorithm maintains the invariant that T is the vertical order of segments intersecting the sweep line.

Time complexity:

- Sorting endpoints: $O(n \log n)$.
- Each endpoint causes $O(1)$ BST ops (insert/delete/find neighbors): $O(\log n)$ each → total $O(n \log n)$.
- Intersection tests are constant time each; only neighbor pairs are tested $O(n)$ times. Thus overall $O(n \log n)$.

(If you must handle segments that share endpoints or collinear overlaps, handle tie-breaking and special cases carefully; but complexity unchanged.)

Example

Segments:

s1: A(1,1) → B(5,5)

s2: C(1,5) → D(5,1)

s3: E(6,6) → F(7,7)

Sweep line moves left → right.

Events sorted by x:

x = 1: A, C

x = 5: B, D

x = 6: E

x = 7: F

Processing at x=1:

Insert s1 and s2. They are neighbors in the sweep-line tree; test intersection → they cross at (3,3). Return "YES".

Hence O(n log n) detects intersection.

Malware / logs problem → decision version is NP-complete

Problem restatement (formal).

- Universe U: all websites.
- For each infected machine i we have subset $S_i \subseteq U$ (sites visited by that machine).
- We seek a smallest collection $H \subseteq U$ such that $H \cap S_i \neq \emptyset$ for every infected machine i. This is the *hitting set* problem (equivalently, the *set cover* dual).

Decision version: Given sets $\{S_1, \dots, S_m\} \subseteq U$ and integer

k

k, is there a set $H \subseteq U$ with $|H| \leq k$ that intersects every S_i ?

Show NP-complete:

1. **In NP.** Certificate is the set H of at most k websites; verify in polynomial time that $H \cap S_i \neq \emptyset$ for all i.
2. **NP-hard.** Standard reduction from Set Cover (or vice versa). For example, Set Cover is NP-complete: given universe U and family $F = \{F_j\}$, ask if there's $\leq k$ sets whose union is U. Convert to hitting set by letting infected machines correspond to elements $u \in U$ and for each

u define $S_u = \{j | u \in F_j\}$ (the indices of sets covering u). A hitting set of size $\leq k$ (choice of indices) corresponds exactly to a set cover of size $\leq k$. This polynomial reduction shows hitting set decision is NP-hard. Hence NP-complete.

Therefore the decision version of “smallest collection of websites visited by all infected computers” is NP-complete.

Example

3 infected computers, their visited websites:

PC1: {A, B, C}

PC2: {B, D}

PC3: {A, D}

Goal: smallest set H that touches all.

Try size $k=1$:

- $H=\{A\}$: covers PC1 & PC3 but not PC2 (which needs B or D)
- $H=\{B\}$: covers PC1 & PC2 but not PC3
- $H=\{D\}$: covers PC2 & PC3 but not PC1
No solution size 1.

Try size $k=2$:

- $H=\{A, B\}$:
 - PC1 visited A/B → OK
 - PC2 visited B → OK
 - PC3 visited A → OK

Thus minimum hitting set = {A, B}, size 2.

This direct maps to a SET COVER instance → NP-complete.

Christofides' algorithm for Metric-TSP and $3/2$ guarantee

Instance: Complete graph $G=(V,E)$ with nonnegative edge weights satisfying triangle inequality.

Christofides' algorithm (steps):

1. Compute a Minimum Spanning Tree (MST) T of G .
2. Let O_O be the set of vertices with odd degree in T . $|O_O|$ is even.

3. Compute a minimum-weight perfect matching M on the induced complete subgraph on O (weights are still metric distances).
4. Form multigraph H by taking edges of $T \cup M$. Every vertex now has even degree $\rightarrow H$ is Eulerian.
5. Find an Eulerian tour of H . Shortcut repeated vertices using triangle inequality to obtain a Hamiltonian cycle (a TSP tour).

Complexity: MST: $O(|E|\log|V|)$. Minimum perfect matching on $|O|$ nodes can be done in polynomial time (blossom algorithm) — polynomial but potentially higher; overall polynomial.

Approximation bound proof sketch ($\leq 3/2$ OPT):

- Let OPT denote length of optimal TSP tour.
- (**MST \leq OPT**): Remove an edge from optimal TSP tour gives a spanning tree of cost $\leq OPT$, so MST cost $c(T) \leq OPT$.
- (**Matching \leq OPT/2**): Consider the optimal tour. Restrict the tour to the vertices in O and take every other edge along the tour to pair up the odd vertices — with triangle inequality this provides a matching whose total cost $\leq OPT/2$. Since M is a minimum perfect matching on O , $c(M) \leq OPT/2$.
- Summing: $c(T) + c(M) \leq OPT + OPT/2 = 3OPT/2$.
- The Eulerian tour of $T \cup M$ has cost $c(T) + c(M)$. Shortcutting (removing repeated vertices) cannot increase cost due to triangle inequality. So the final TSP tour cost $\leq 3/2$ OPT.

Thus Christofides returns a tour with cost at most $3/2$ times optimal.

Example

Complete graph with vertices $\{A, B, C, D\}$ and metric distances:

	A	B	C	D
A	0	4	6	7
B	4	0	3	5
C	6	3	0	2
D	7	5	2	0

Step 1 — MST

Pick minimum edges:

C-D (2), B-C (3), A-B (4)

MST cost = $2 + 3 + 4 = 9$

Odd-degree vertices = $\{A, D\}$

Step 2 — Perfect Matching

Match A–D (cost 7)

Total = 9 + 7 = 16

Step 3 — Euler Tour & Shortcut

Euler tour: A → B → C → D → A

No shortcuts needed.

Cost = 16

Optimal TSP = 14

So Christofides = $16 \leq 1.5 \times 14 = 21$

(Guarantee holds).

Approximation algorithms for Vertex Cover and Set Cover

Vertex Cover — 2-approximation (simple greedy / matching based)

Algorithm (edge-picking):

- Initialize cover $C = \emptyset$.
- While there is an uncovered edge (u,v) :
 - Add both endpoints u and v to C .
 - Remove all edges incident to u or v .
- Return C

Analysis:

- Each iteration picks an edge and forces adding 2 vertices; set of picked edges forms a matching M (no two picked edges share a vertex). Any vertex cover must pick at least one endpoint from each matching edge, so $|OPT| \geq |M|$. Our cover size $|C| = 2|M| \leq 2|OPT|$. So 2-approximation.
- Time: $O(m)$ if implemented carefully.

Alternative: LP relaxation + rounding or maximal matching give same 2-approx ratio.

Set Cover — greedy algorithm ($H_n \approx \ln n$ approximation)

Problem: Universe U of n elements, family F of subsets. Pick $\leq k$ sets to cover U with minimum number of sets.

Greedy algorithm:

- Start with $C = \emptyset$, uncovered elements $R = U$.
- While $R \neq \emptyset$: pick the set $S \in F$ that covers the largest number of uncovered elements $|S \cap R|$. Add S to C and remove covered elements from R .
- Return C .

Approximation factor and proof idea:

- Greedy achieves a H_n approximation where $H_n = 1 + 2 + 3 + \dots + n \leq ln(n+1)$.
- Proof (standard): Let OPT be size of optimal cover. At each greedy step, the chosen set covers at least $1/OPT$ fraction of remaining uncovered elements. This yields a geometric decay in uncovered elements; summing gives $\leq H_n \cdot OPT$ sets chosen.

Complexity: naive implementation $O(nm)$ or with clever book keeping $O(m \log m + \text{total set sizes})$.

Other methods: Primal-dual or LP rounding give same $O(\ln n)$ guarantee; hardness results show $\Omega(\log n)$ approximation lower bound (unless $P=NP$), so greedy is essentially optimal.

A. Vertex Cover 2-Approx Example

Graph edges: (A,B), (B,C), (C,D), (D,E)

Greedy algorithm picks an uncovered edge each time:

Pick (A,B) → add A,B

Pick (C,D) → add C,D

Now E is isolated → done

Cover = {A, B, C, D} (size 4)

But optimal = {B, C, D} (size 3)

Ratio = $4/3 \leq 2 \rightarrow$ valid 2-approx.

B. Set Cover Greedy Example

Universe $U = \{1, 2, 3, 4, 5\}$

Sets:

$S_1 = \{1, 2\}$

$S_2 = \{2, 3, 4\}$

$S_3 = \{4, 5\}$

$S_4 = \{1, 5\}$

Greedy picks:

1. S_2 (covers 3 elements: 2, 3, 4)
Remaining: {1, 5}
2. S_4 (covers {1, 5})

Total sets used: 2

Optimal is also 2, so greedy = optimal in this example.